

Performance Comparison of Several Neural Networks in Policy Approximation in a Ludo Game Playing Agent

Kartik Tushir
kk1478@msstate.edu

Rajeev Jogi
rj993@msstate.edu

James Kastrantas
jgk16@msstate.edu

Index Terms—reinforcement learning, policy approximation, ludo, deep learning, gradient ascent, experience replay

I. TASK DEFINITION

Our task for this project is to determine which neural network models work best for a reinforcement learning agent playing the game of Ludo. A recurrent neural network (RNN), convolutional neural network (CNN), and fully-connected artificial neural network (ANN) will be used to approximate a policy for the game.




Fig. 1. A Ludo game board labeled with encodings for pawn positions. [1]

Ludo is a game in which each player tries to move their four pawns from the starting position to the goal in the center of the board. Player can move one of their pawns from one to six spaces decided by a six-sided die roll. A player can remove another player's pawn by moving into the opposing pawn's space, but there are areas of the board where opposing players cannot be reached or are immune to removal. The action available to a player is the decision of which pawn to move. This decision must be weighed against the current state of the game, making some pawns better to move than others. The agents must be trained to learn these strategies.

II. LITERATURE REVIEW

A. Reinforcement Learning: An Introduction [2]

A more in-depth discussion of policy approximation is given in the section on policy gradient methods. Advantages of policy approximation are given, but missing are discussions on the disadvantages to policy approximation. Converging on a policy approximation by computing the Monte-Carlo policy gradient will be essential for estimating a policy, and that is outlined here.

B. Reinforcement Learning for Robots Using Neural Networks [3]

A dissertation that provides techniques for effective reinforcement learning. The background provides an excellent introduction to reinforcement learning. Discussed is hierarchical learning, using smaller elementary problems that can agent can solve which can be used to learn strategies for more complex problems. Experience from human teachers to improve training is discussed. Experience replay, which is essential for assigning rewards to state-action pairs, is also discussed. All of these topics may be useful to training our agents.

C. Q-Learning [4]

This technical note proves that Q-learning converges on an optimal policy. It serves as a short introduction on Q-learning, and elaborates on the relationship between Q-learning and dynamic programming. It also shows that each problem instance may have multiple optimal policies, but each state-action pair (s, a) has a unique Q-value $Q(s, a)$. Using the state-action pair that maximizes the Q-value is the same as having an optimal policy π^* . This is important because the reward table \mathcal{R} and transition probability table P_{xy} may not be known ahead of time, or is too large to compute beforehand, and using Q-learning eliminates the need for these resource-intensive models.

III. MOTIVATION

It is in our experience that teaching concepts to others is the best way to learn those concepts. The motivation for the project comes from the requirements for the class, but the final result will be shaped as an assignment appropriate for

an introductory, undergraduate level AI class, complete with implementation for them to get started with.

In a two-player game of Ludo, there are no more than 2 quadrillion possible states of the game. While not impossible, utilizing immense amounts of secondary memory is necessary to compute a policy table. A reinforcement learning agent that can generalize from experience and approximate this table using a neural network would be beneficial for both people training the agent and people running the agent on limited computational resources such as personal computers or smartphones.

Simply explaining to students how these agents work does not provide a good hands-on experience for students learning reinforcement learning. Forming an assignment that gives them a completely-observable, zero-sum game with a large state space that makes it very, very difficult to compute and store a policy table for the game, we can give them the experience they need to design their own reinforcement learning agents. This can lead into discussion on larger problems complicated by partial-observability.

IV. METHODS

Adapting LUDOpY from the GitHub repository as the implementation of the environment, we are going to train several reinforcement learning agents to learn and compete at a 2-player game of Ludo. Although Ludo can be played with 4 players, 2 players will be used in training the agents to simplify the state space.

Episodes generated by LUDOpY's random walks will allow the agents to assign rewards to states based on their future outcomes. Gradient ascent will be used to adjust the weights of the policy network since this is a maximization problem.

Several neural network models for policy approximation: RNN, CNN, and fully-connected ANN. This will allow students working in groups to distribute training tasks evenly, and allow them to compare game-playing performance of their agents.

V. EVALUATION

We will see which networks best maximize future rewards by measuring future rewards from the initial state. We can also setup brackets and play games against the other agents and determine which agent wins more games. Parameter size is also a factor we will consider. We can also look at adjustments to the hyperparameters of the network to see which models converge faster according to a different learning rate or depth of the network.

REFERENCES

- [1] SimonLBSørensen, LUDOpY, GitHub repository, <https://github.com/SimonLBSørensen/LUDOpY>, 2020.
- [2] R. S. Sutton and A. G. Barto, Reinforcement Learning: An Introduction, 2nd ed. Cambridge: MIT Press, 2020. Accessed on Sep. 26, 2020. [Online]. Available: <http://incompleteideas.net/sutton/book/RLbook2020.pdf>
- [3] Long-Ji Lin, "Reinforcement learning for robots using neural networks," Ph. D dissertation, School of Computer Science, Carnegie Mellon Univ., Pittsburgh, 1993. Accessed on Sep. 11, 2020. [Online]. Available: <http://isl.anthropomatik.kit.edu/pdf/Lin1993.pdf>
- [4] C. J. C. H. Watkins, P. Dayan, "Q-Learning," Machine Learning, vol. 8, pp. 279-292, May 1992.